



University
of Glasgow | School of
Computing Science

Design and implement your own programming language

Kyle Simpson
Kristiyan Dimitrov
Darren Findlay
David Creigh
Gerard Docherty

Level 3 Project — 27 March 2015

Abstract

The abstract goes here

Education Use Consent

We hereby give our permission for this project to be shown to other University of Glasgow students and to be distributed in an electronic format. **Please note that you are under no obligation to sign this declaration, but doing so would help future students.**

Name: _____ Signature: _____

Name: _____ Signature: _____

Name: _____ Signature: _____

Name: _____ Signature: _____

Name: _____ Signature: _____

Name: _____ Signature: _____

Contents

1	Introduction	4
1.1	Project Outline	4
1.2	Motivation	4
1.3	Aim	4
1.4	WebGL Background	5
1.5	Outline of report	5
2	Language Tutorial	6
2.1	Getting Started	6
2.2	Variables and Arithmetic operators	7
2.3	If statement and comparative operators	9
2.4	Modifying Operators	10
2.5	Functions and Scope	10
3	Language Reference Manual	12
4	Project Plan	13
5	Language Evolution	14
6	Compiler Architecture	15
6.1	Overview	15
6.2	Lexer and Parser	16

6.2.1	Introduction	16
6.2.2	Initial steps	16
6.2.3	Design	17
6.3	Code Generator	19
6.4	Abstract Machine	19
6.5	Shader Manager	19
6.5.1	Introduction	19
6.5.2	Definitions	20
6.5.3	Requirements	20
6.5.4	Design Process and Issues	20
6.5.5	Usage	21
6.5.6	Future Changes and Evaluation	23
6.6	Module Facade	24
7	Development Environment	25
8	Online IDE	26
8.1	Design Process	26
8.2	Implementation	27
8.3	Evaluation	28
9	Test Plan and test suites	29
10	Conclusion	30
10.1	Contributions	30
10.1.1	Report	30
10.1.2	Program	30
A	Appendix A: Source Listing	31

Chapter 1

Introduction

1.1 Project Outline

For our project we chose project number 3924; design and implement your own programming language. This project involves designing a programming language, by creating a BNF, system diagram, and the syntax and semantics of the language, among others and then implementing it, using a parser, compiler, and abstract machine.

1.2 Motivation

Our Motivation behind the choosing of this particular project is that, especially in the past few years, computing has come to the forefront of teaching. Many parents now want their children to learn how to code to some degree, with computing science now being a much more prevalent subject in primary and secondary schools. Now in our third year of university, we have all taken part in a fair amount of coding, however, none of us have ever attempted to write our own language to code in. This project seemed challenging, but achievable, and very rewarding not to mention interesting.

1.3 Aim

The aim for this project is to create a graphical programming language that is simple to use and read, and can be picked up by nearly any novice user, but is also powerful, and can be used to create drawings, or edit pictures. Our idea for a simple language came from Processing, a graphical language created by Casey Reas, and the basis for the language is going to be WebGL a browser version of OpenGL. With this project, we want to create a language that follows what we think are useful and intuitive features in the design of most programming languages, for example using `for` functions and loops and keeping the familiar structure of the `for` and `while` loops, and making sure our language follows them, while finding what we believe to be confusing or annoying common features of programming languages, like the double equals for the comparison of two objects, and replace them with what we find suitable and more

1.4 WebGL Background

Web Graphics Library (WebGL) is a JavaScript API, that makes use of the `<canvas>` element available in HTML to render interactive 2d/3d graphics. It is a branch of the Open Graphics Library Embedded Systems 2.0 (OpenGL ES 2.0), and is supported by most modern browsers, like Google Chrome and Mozilla Firefox, with no necessary plugin (as it makes use of `<canvas>`). A WebGL program consists of control code written in JavaScript, and function code (shader code) that is executed by a computer's GPU.

1.5 Outline of report

The remaining sections of this report will cover the following topics:

- Language Tutorial :- This will comprise of a brief introduction to our programming language, which we have named **Sketch**
- Language Reference Manual :- This will contain the grammar, and describe the language
- Project Plan :- This will cover a timeline and basic structure of our project, and what we aimed to achieve from it. **I think so anyway**
- Language Evolution :- This will describe the decisions we made for the language syntax and domain, and why we made them.
- Compiler Architecture :- This section will cover the 5 different sections of the compiler:
 - Lexer and Parser
 - Code Generator
 - Shader Manager
 - Abstract Machine
 - Module Facade
- Development Environment:- The environment(s) that we used to help structure and create this programming language
- Test Plan and Test Suites :- An overview on how we tested the language, using what, and any bugs we discovered
- Conclusions :- This will cover what, as a team, we learned and discovered during this project, any problems we encountered, and anything we would try differently.

Chapter 2

Language Tutorial

This chapter comprises of a brief introduction of how to use our language.

2.1 Getting Started

In every language, the first program to write is always 'hello world' - where you would print the words "hello world". However, as this language is mainly a graphical language, the equivalent would be 'hello square', and the aim is to draw a square. In Sketch, the program required to draw a square is:

```
function init() {  
  Polygon square = ({0,0} + {5,0}) * 4;  
  draw square;  
}
```

Listing 2.1: Hello Square

To compile any Sketch program, insert the code into the browser page for our project and click "Render". The result of the program will then be shown on the canvas.

Any Sketch program you write will have to consist of at least one function, and variables. These functions and variables can be named anything you like. You can call other functions to help carry out the task, only ones that you have written, as there are no libraries provided in the current version of the language.

In Sketch, implicit semi-colons exist at the end of every line. This means that, you don't have to end each line with a semi-colon, and whether you do or not will not have an effect on the compilation of the program.

The draw() function is used to pass the shapes you want to draw to the correct function, so it can be written to the canvas. This will be discussed in detail later.

One way in which you can pass data between functions is by including variables in the calling state-

ment as arguments. However, you can only do this if the function that is being called is expecting the same number and type of variables. An example is:

```
function init() {  
    int n = 2;  
    function takesParams(int n) {  
        print(n);  
    }  
    takesParams(n);  
}
```

Listing 2.2: Function Parameters

Here, you see that the calling statement - *takesParams(n)*; - provides the correct number and type of arguments. If there was another argument included, for example *takesParams(n, 7)*, or the wrong types, then it would not compile. This will be talked about in more detail in Chapter 2.5. As you can see in Listing 2.1, the code inside of a function is enclosed by curly braces {}. The draw keyword is used to draw any shapes that you have constructed to the canvas. It will take one parameter, either the point, line or polygon that you have constructed. This is a predefined keyword in the language, so you don't have to write it, but you cannot name any of your functions *draw()* as a result.

2.2 Variables and Arithmetic operators

This next section will use a more complicated program than before, introducing more features, such as comments, loops, and expand on previously touched on features, such as variables.

```
#This Program will draw 3 different shapes - triangle, square and pentagon.  
//Both of these lines are ignored by the compiler.  
  
function init(){  
    int sides;  
    sides = 5;  
    int counter = 3;  
  
    Point pt1 = {3,1}  
    pt2 = {1,3}  
  
    Line l1 = (pt1 + pt2);  
  
    #this will loop for the number of sides  
    while(counter <= sides){  
        clear;  
        Polygon shape = l1 * counter;  
        draw shape ;  
        counter++;  
    }  
}
```

Listing 2.3: Comments and Point Addition

The initial lines indicate how to show comments in your code in Sketch. Using a hashsign (#) or a double slash (//) will cause the compiler to ignore all characters until the next newline, and show that the current line is a comment. These can be used to explain how your programs works, and make it easier to read and understand, for you or other users. The 'clear' function in this example

is a reserved keyword, and is used to clear the canvas. A reserved keyword is a word that cannot be used otherwise, for variable names, of function calls - 'draw' is one such example we have shown.

As you can see in Listing 2.3, there are two ways of declaring variables, explicitly and implicitly. Each variable declaration must include the variable name, which is an alphanumeric string that always begins with a letter. To explicitly define your variable, each variable name must be preceded by the intended type. For example, the variables 'sides' and 'counter' in Listing 2.3 are declared explicitly. Explicit declaration allows for the ability to not initialise the variable immediately. Implicit declaration is when the type is not defined by the user, but instead taken from the context. The type is inferred from the initialization of the variable afterwards. For example, 'pt2' in Listing 2.3 is implicitly defined, and from the context, the compiler will recognise that it is of type Point. However, with implicit types, you must immediately initialise the variable so its type can be inferred.

The while loop in Sketch is of standard format and standard functionality. This means that the condition contained in the parentheses is tested, and if it returns true, then the body of the loop, and whatever statements it may contain, are executed. Then the original condition is retested. If true, it again executes the body. This continues until the condition is false, at which point, the body of the loop is not executed, and the next command outside the loop is executed and the program continues. A while loop has the format:

```
while(condition) {  
    #insert body here  
}
```

Listing 2.4:
The "While"
Loop

Sketch has three main types for drawing shapes, lines and so forth. These are Point, Line and Polygon. Let's start with Point; Point is the basis of all possible shapes. Points make the vertices on which the shapes will be drawn about. A point is defined in the following format {*x* Coordinate, *y* Coordinate}. An example of both implicit and explicit declaration is shown in Listing 2.3 with 'pt1' and 'pt2'. These points can be added together to create a Line. Lines are a collection of two points joined together. The format for defining a line, shown by 'l1' in Listing 2.3, is Line = (Point1 + Point2). Lines can be extended from the centre point by multiplying it by a number. For example, multiplying a line by 4 will make it 4 times the size. Finally, variables of type Polygon are a collection of lines, displayed as a closed circuit. Polygons can be defined by either adding multiple lines together, but they must all join, or by defining one line, and multiplying it by the number of sides you would like. This would create a shape of the stated number of lines, closed off, with the origin being the first co-ordinate of the line. The format of this would be either Polygon = (Line2 + Line2 ... + LineN) or Polygon = Line * N. In these examples, N is the number of sides desired.

The line 'counter++' increments the counter, using the modifying operator '++'. We will explain modifying operators later in this tutorial.

Sketch also has one more type of loop - the "for" loop. The for loop has a completely different structure to the while loop, but has, more or less, the same functionality. The choice between these two types is based on the context, which one would make more sense, for example, incrementations should be used in for loops, while boolean variables should be the control in while loops. The following is the while loop from Listing 2.3, changed to a for loop:

```

for(int counter = 3; counter < 5; counter ++) {
    ...
}

```

Listing 2.5: The "For" Loop

The for loop statement has three parts. Firstly, the control variable is initialised (it also does not need to be declared beforehand), then the statement that will be checked after every execute of the body, and finally, the incrementation of the control variable. The general functionality is the same as the while loop.

2.3 If statement and comparative operators

This section will cover how to perform comparing statements in Sketch. The following is an if statement which compares to check if a variable is equal to something else:

```

function init() {
    int n = 5;
    if(n == 5) {
        print(n);
    }
    else if(n > 5) {
        print("less than 5");
    }
    else if(n < 5) {
        print("greater than 5");
    }
}

```

Listing 2.6: If-Else Conditionals

If statements are a way of executing a series of commands, but only if a particular statement returns true. To compare to elements, a comparative operator is needed. In Sketch, there exists a series of six comparative operators:

- `==` will return true if left hand side is equal to right hand side.
- `>` will return true if left hand side is greater than right hand side.
- `<` will return true if left hand side is less than right hand side.
- `!=` will return true if left hand side does not equal right hand side.
- `!>` will return true if left hand side is not greater than right hand side (ie is less than or equal to).
- `!<` will return true if left hand side is not less than right hand side (ie is greater than or equal to).

An exclamation mark can be inserted in front of the check statement in the if, to receive the inverse of what it returns.

2.4 Modifying Operators

This section will cover how you can make use of sketch's modifying operators. The following example will outline them:

```
function init() {  
  int n = 1; #line 1  
  n += 2; #line 2  
  n *= 2; #line 3  
  n /= 2; #line 4  
  n -= 2; #line 5  
  n %= 2; #line 6  
  n++; #line 7  
  n--; #line 8  
  n = (n + 3) * 5; #line 9  
}
```

Listing 2.7: Modifying Operators in Action

Listing 2.7 outlines many modifying operators in Sketch. To explain each one, we will walk through the program itself. As previously mentioned, the first line will declare `n` to be of type `int`, with a value of 1. The next line will change the value of `n` by adding one to it. The general syntax would be **variable to be changed* += *how many to change it by**. The next three lines follow the same syntax and semantics, with the exception that, line 3 would multiply the variable by the second value and assign it, line 4 would divide it by the second value and assign it, and line 5 would subtract by the second value and assign it, and finally line 6 would change the variable to equal the value of it modulo the second variable. This means that it would change it to the remainder after being divided by that number. So, after line 2, `n` would be equal to 3, after line 3, `n` would equal 6, after line 4, `n` would equal 3, after line 5 `n` would equal 1 and after line 6, `n` would still equal 1, as the remainder of $1 / 2$ is still 1. These operators are functionally a shorter way of writing `n = n + *value*`, but obviously changed respective to the context (division, multiplication, subtraction and modulo).

Line 6 is different. This operator will increment the variable by only 1. In the context of this program, `n` will now equal 2. The following line following the same structure, but instead subtracts one, meaning `n` is back to equal 1. These can be useful in loops for example, to increment or decrement your controlling variable as necessary. The final line outlines how parenthesis work in regards to performing mathematic actions. The flow it would take, is to first, add to `n`, making it equal 4, and then multiply by 5, hence `n` equals 20. The brackets would specify what task has to be carried out first, as if they were not here on this occasion, by the laws of maths, 3 would be multiplied by 5 and then added to `n`, giving a different outcome. All of these modifying operators, except on line 9, are functionally just shorthand, but can be very useful and time saving while writing a program.

2.5 Functions and Scope

In Sketch, functions, as mentioned before, taken certain types and number of parameters dependant on what you specify, it can range from 0 to `n`. Furthermore, these functions can be called whatever you decide, and can either have a return type or not, dependant on your needs. Example 1.6 will show these in action:

```

function init(){
    int x = 1;
    int i = 5;
    function foo(int n) -> int{
        n++;
        return n;
    }
    function foo2(int n){
        n++;
    }
    i = foo(int i);
    foo2(int x);
}

```

Listing 2.8: Functions and Scope

In Listing 2.8, you can see different properties of defining and calling a function. If a function is not returning anything, then you can see that the void is implicit. Furthermore, if a function has a return type, then you can assign a variable to equal its returned value, if it is of the same type. In regards to scope, functions have local scope, meaning that they would alter outside variables, or variables in other functions. This means that, as shown here, you can call variables the same name in two different functions, and there will be no crossover or confusion of values or any conflicts.

Chapter 3

Language Reference Manual

Chapter 4

Project Plan

Chapter 5

Language Evolution

Chapter 6

Compiler Architecture

6.1 Overview

Our program module, as part of the role it plays, contains a combination of a compiler, bytecode executor and WebGL interface - written in JavaScript as required by the platform. This section aims to provide a view of the overall architecture and connection between these components, as well a brief explanation of each component's function and design considerations.

The core language module can be divided into 5 main parts:

1. **Lexer and Parser** — The program component charged with tokenising and interpreting Sketch programs, generating an abstract syntax tree for use in code generation.
2. **Code Generator** — The program component responsible for conversion of abstract syntax trees into bytecode sequences, for execution by the abstract machine.
3. **Shader Manager** — An active subsystem of the module responsible for management, selection and execution of shader programs for drawing at runtime.
4. **Abstract Machine** — The main active subsystem present in the module, responsible for all runtime code execution. The abstract machine relies on the shader manager to make draw calls, but handles direct WebGL manipulation for certain key functions separately.
5. **Module Facade** — The active component responsible for managing the interconnections and operation of the above components, as well as exposing the module interface to programmers.

Individual and detailed discussion of each system is provided below. The module's subsystems are then connected as shown in Figure 6.1.

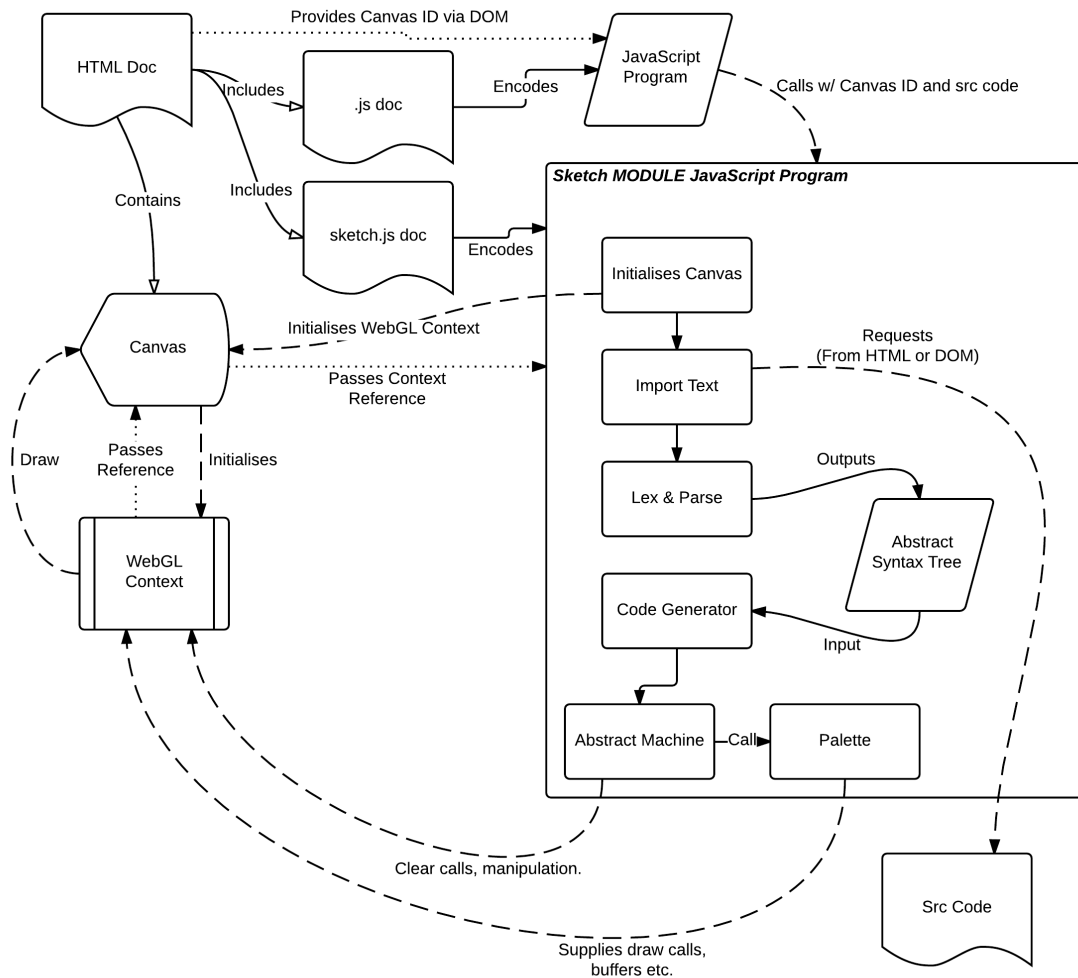


Figure 6.1: Overall system diagram of Sketch, including data flow and interactions with the browser.

6.2 Lexer and Parser

6.2.1 Introduction

The lexer and the parser are both part of the syntactic analysis of the compiler. During this stage, a source program is parsed in order to determine if it is well-formed, and to determine its phrase structure, in accordance with the syntax of our programming language. If the source program complies with the syntax and grammar of Sketch, an abstract syntax tree (AST) will be returned from the parser. An AST is a way to represent a source program's phrase structure.

6.2.2 Initial steps

At the beginning of implementing parser and lexer for our grammar, an appropriate parser generator for JavaScript had to be found. We found three possible parser generators, each of which would

helped us with desired result. ANTLR, JISON and PEG.js were considered as possibilities.

ANTLR uses LL(*) parsing and it can generate lexers, parsers, tree parsers and lexer-parsers. Parsers produced by it are able to automatically generate an AST, which was necessary for our design. However, the most recent release of the software (ANTLR 4) does not support JavaScript as a target source, even though the previous version does.

JISON and PEG.js do not encounter the same problem - they successfully support JavaScript as a target and both could have been appropriate choice for our parser. PEG.js is based on parsing expression grammar formalism, which is more powerful than traditional LL(k) and LR(k) parsers. On the other hand, Jison can recognize languages described by LALR(1) grammars and provides additional features (i.e. the ability to define operator associations and precedence). Because of this additional features, JISON was our choice for parser generator.

6.2.3 Design

JISON provided us with the opportunity to specify both tokenizing rules and language grammar in the same file. This turned out to be very helpful in latter stages when the communication between the parser and code generator had to happen. The whole design process is comprised of three stages: lexical analysis, specification of the grammar and specification of the abstract syntax tree in JSON format.

6.2.3.1 Lexer

Lexer is a program or function that performs lexical analysis (the process of converting a sequence of characters into a sequence of tokens). A token consists of a token name and attribute value. The token name represents a kind of lexical unit (i.e. keyword, identifier) and later the token names are processed by the parser.

For each keyword of our language, a token was specified in the grammar file. Each of the operators and punctuation symbols has a specified token as well. There are tokens for the identifiers, numbers and strings as well. This way, all the possible tokens for Sketch were covered in the grammar file. In Listing 6.1, an example for each of them is presented.

```
%lex
%%
...
"if"                return 'IF';
"Line"              return 'LINE';
...
"{"                 return 'OPEN_BRACE';
...
"+"                 return 'PLUS';
...
[0-9]+ ("." [0-9] *)? return 'NUMBER';
[a-zA-Z_]+ [a-zA-Z0-9_]* return 'IDENTIFIER';
/lex
```

Listing 6.1: Tokens of Sketch

6.2.3.2 Grammar

Next stage in the implementation for this part of the compiler was to specify the grammar rules for our graphic programming language. All the rules specified in the initial BNF were translated in a format that the parser generator can accept.

As a result, in the case when a given rule of the grammar can be used, that rule had to be rewritten again but this time using the recursion feature of JISON. JISON supports 2 types of recursion - left and right, but our grammar uses left recursion, since it more reliable and it can parse a sequence of any number of elements with bounded stack space.

Another feature of JISON became very useful in the elimination of shift-reduce conflicts - the ability to specify precedence. Initially our grammar was full of shift-reduce conflict because of the "dangling else" in "condition_statement" (Listing 6.2). By using precedence, our grammar become less ambiguous and errors are less likely to occur.

```
condition_statements
: IF OPEN_PARENS exp CLOSE_PARENS statement
                                %prec IF_WITHOUT_ELSE
  ...
| IF OPEN_PARENS exp CLOSE_PARENS statement ELSE statement
  ...
;
```

Listing 6.2: condition_statements in Sketch after removal of Shift-Reduce conflict

6.2.3.3 Abstract Syntax Tree (AST)

If a source program in Sketch uses the correct syntax and grammar rules, after parsing an Abstract Syntax Tree would be generated and passed to the code generator. Our parser produces a tree represented as an object in JSON format. As a result, the JISON file had to be changed and for each grammar rule the corresponding JSON object was added. In Listing 6.3, the output for "Hello Square" program is show. If one of the grammar rules is violated, no AST will be generated and error will be thrown - this feature is provided for us by the parser generator.

```
1 {
2   "type": "function",
3   "arguments":
4     [ "init", "", "void",
5       [ { "type": "variable-decl-assign",
6           "arguments":
7             [ [ "Polygon", "square",
8                 { "type": "multiplication",
9                   "arguments":
10                     [ { "type": "addition",
11                       "arguments": [ [ "0", "0", [ "5", "0" ] ] ], "4"
12                     ]
13                   } ] ] },
14             [ "draw", "square" ]
15           ]
16         ]
17     }
```

6.3 Code Generator

6.4 Abstract Machine

6.5 Shader Manager

6.5.1 Introduction

In designing our system, we established that regardless of the chosen compilation and execution pathway all of our library's shaders would have to be intelligently managed; both during runtime and for storage purposes inside the module itself. While an extensive system like Palette may seem too broad in scope for Sketch's module in its current form, at the language's inception we wanted custom shapes and shaders to be a defining feature of our platform. As such, Palette allows us to one day realise our initial ambitious vision of Sketch due to the level of extensibility and portability it provides.

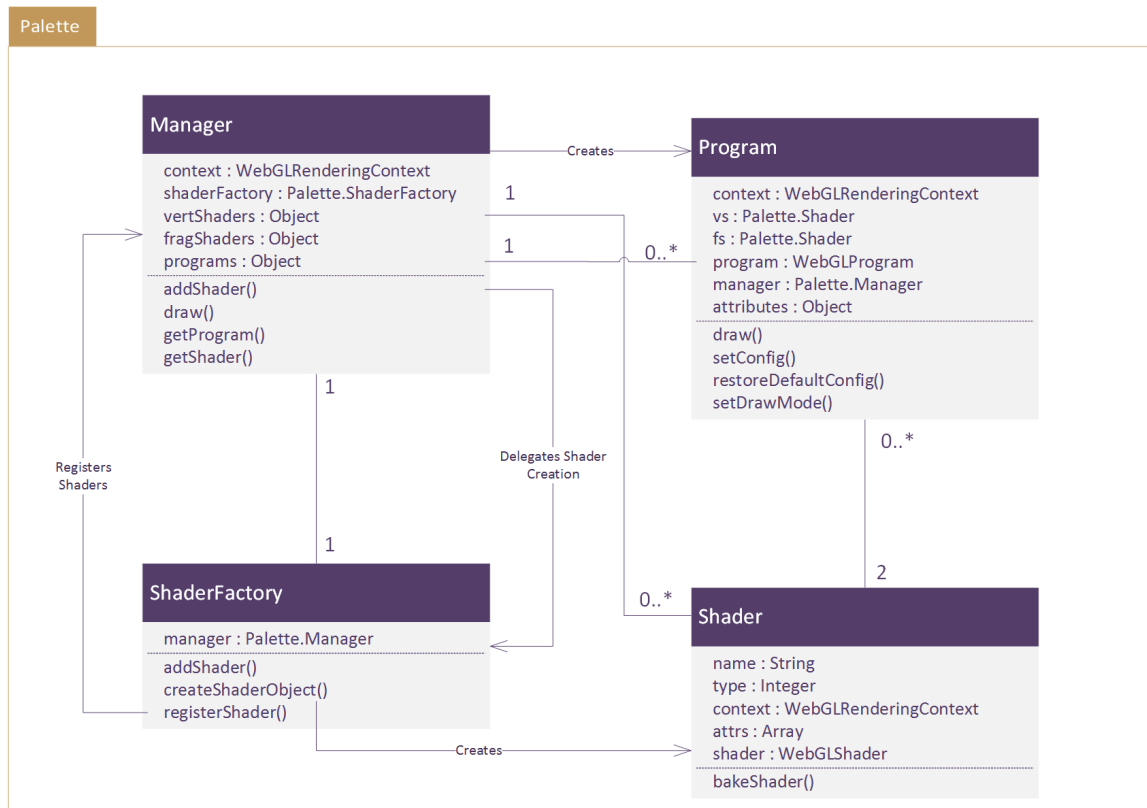


Figure 6.2: A class diagram of Palette.

6.5.2 Definitions

Shaders are the graphics card level program units used in the conversion from vertex arrays into two-dimensional projections in the screen space, and in the selection of pixel colours to populate the screen space from that projection. In WebGL, older versions of OpenGL and in its subsets such as GLES, shaders must always come in pairs: a *vertex shader* and *fragment shader* comprise each *shader program*. This may then be accessed, modified and called through the *WebGL API*.

6.5.3 Requirements

The unique needs of our module created an interesting set of design requirements for the subsystem - mandating a somewhat in-depth understanding of the WebGL API to manipulate the canvas as required. Upon examination, the shader manager's requirements were found to be:

- **Extensibility** — Sketch needed to support future extension with various shapes, shaders and other such additions beyond its standard library - most shader programs require writing specially tailored code in an object's `.draw()` method. Similarly, modifications to shader code often require changes at the software level if the attributes or buffers must be changed - this coupling must be minimised or eliminated.
- **Portability** — Such tailored draw code cannot be imported without putting users at significant security risk by requiring them to execute arbitrary code. Some way to infer the attribute update process from the shader itself is a necessity.
- **Ease of Use** — Definition of new shader files should be intuitive and an easy process, provided the user has valid shader code. The linking and compilation of shaders and shader programs must be handled internally, and completely sequestered from external control to centralise most of the API manipulation.
- **Configurability** — Shader programs should have variable control methods exposed as an external interface, to provide a simple means of control over the intricacies of the WebGL API. For instance, programmers must be able to change the draw mode on a moment's notice, or even extract the shader program pointer if they require direct access to WebGL.
- **Simplicity** — Interactions with the module should be designed with brevity in mind - being able to execute variable update and draw functions as a single method call makes WebGL usage simpler in the other subsystem implementations - lowering the barrier of entry. WebGL is plagued by esoteric function names, inconsistent function prototypes and an arcane set of named constants: all of which must be used regularly to render graphics.

6.5.4 Design Process and Issues

The first challenge encountered stemmed from trying to find a way to pass shape definitions within our language. At first, we reasoned that since shapes and objects traditionally have predefined draw functions then we'd need to find a way to export raw JavaScript for each shape class. As mentioned in the requirements, we found that direct code evaluation would not only require difficult scoping

considerations but would also present a massive security risk to users of the system. From here, we considered merely assigning shape classes a pair of shaders - this seemed to be the logical choice.

We had originally considered allowing programmers to add shader object JSON from a URL source; this would provide an even more convenient set-up for the manager by reducing shader registration to a single line of code per file. However, we ran into trouble with the management of asynchronous execution calls caused by the main thread running while the download completed. We briefly considered making the HTTP request synchronous, but this caused the entire window to freeze until downloading had completed - we quickly discarded this plan for the time being.

Although our JSON system of "[name, type, default]" worked well for numeric and matrix attributes, we were originally at a loss for how buffers and vertex attributes should be handled. By forcing programmers to supply a buffer object with a known name as well as a vertex attribute specific to their shader code, we managed to bypass this problem with regards to the vertex buffer. We found that using the default value to store the name of the referenced buffer was an eloquent solution - keeping each array to size 3. As far as buffers were concerned, we then realised that we could perform the same substitution to store the itemSize of a buffer. There are many other possible configuration options which had to be cast aside (stride, offset, etc.) to allow this simplicity, however.

Midway through the development of the abstract machine, we realised that we needed a means to convert from planar polygons (like those that Sketch works with) into triangle sets to render any geometry to the screen. While researching polygon triangulation algorithms we came across *earcut* (<https://github.com/mapbox/earcut>), a JavaScript library which provides the features and speed that our platform would need to make use of. By patching in an extra draw mode to Palette on top of the standard WebGL modes, the simplicity of execution at the abstract machine level was maintained.

While discussion of the implementation details themselves would be appropriate in a dedicated paper, they must unfortunately be omitted as they are outside of this report's scope.

6.5.5 Usage

While JSDoc automatically generates the API documentation for Palette, construction of Shader JSON files is not yet documented. As these are an essential part of Sketch's operation, and shader source code will not suffice on its own, a light explanation is provided.

Shader JSON objects may fall into three types:

1. **Vertex Shaders** — Objects following the form of Listing 6.4, with type set to 0 and at least two attributes: a *buffer* named "vertexBuffer" with default value set to the number of components per vertex, and a *vertexAttrib* with a name matching an attribute within the shader and default value set to "vertexBuffer". Further attributes are defined in the same format as fragment shaders.
2. **Fragment Shaders** — Objects defined similarly to Listing 6.5, with type set to 1 and no restriction on what attributes must be predefined.
3. **Shader Lists** — Objects following the structure of Listing 6.6, with type set to 2 and the content tag corresponding to an array of further shader objects of any type.

All integer and vector types are supported, along with their array equivalents. These are denoted as *int*, *int[]*, *float*, *float[]* with vectors taking the form *ivec_n*, *ivec_n[]*, *vec_n* and *vec_n[]* for $n \in [2, 4] \cap \mathbb{N}$. Float matrices may be defined by *mat_n*, for n within the same domain.

```

1 {
2   "type": 0,
3   "name": "squareVtx",
4   "src": "#version 100; attribute vec3 aVertexAttrib;//...",
5   "attrs": [
6     ["aVertexAttrib",
7      "vertexAttrib",
8      "vertexBuffer"],
9
10    ["vertexBuffer",
11     "buffer",
12     3]
13  ]
14 }
```

Listing 6.4: Vertex Shader JSON Structure

```

1 {
2   "type": 1,
3   "name": "squareFrag",
4   "src": "//source",
5   "attrs": [
6     ["name",
7      "type",
8      "default"]
9  ]
10 }
```

Listing 6.5: Fragment Shader JSON Structure

```

1 {
2   "type": 2,
3   "content": [
4     {
5       "type": 0,
6       "name": "...",
7       "src": "...",
8       "attrs": [...]
9     },
10    {
11      "type": 1,
12      "name": "...",
13      "src": "...",
14      "attrs": [...]
15    }
16  ]
17 }
```

Listing 6.6: Shader List JSON Structure

Initialisation of Palette as part of a JavaScript program is a very simple process. Once the module has been included in the execution space through either loading by the browser over HTTP or concatenation as part of a build script, the API is accessible. An example integration is as follows:


```

1 //Given a <canvas> element with id "myCanvas".
2 //The shader text has been imported either by XMLHttpRequest
3 //Or has been embedded as a string literal into variable 'text'
4
5 //Retrieve elements from DOM, and the context from the canvas.
6 var canvas = document.getElementById("myCanvas");
7 var glctx = canvas.getContext("webgl");
8
9 //Create and set up Palette.
10 var manager = new Palette.Manager(glctx);
11 manager.addShader(text);
12
13 //Now we're ready to draw!
14 manager.draw("vsName", "fsName", [/*DATA*/],
15             { /*vsConfig*/ }, { /*fsConfig*/ });

```

Listing 6.7: Palette Usage Example

The full API documentation is available under docs/index.html within our repository (and the deliverables folder supplied) under the namespace "Palette".

6.5.6 Future Changes and Evaluation

While Palette provides a sizeable kernel of features, there is still much room to improve the module - not only to supply all of the functionality we'd originally planned, but also to enable its use in external projects of all scopes.

- **Texture Support** — an absolute necessity to enable images and texturing within Sketch, and for almost all conceivable WebGL projects. Textures require vast amounts of configuration - the issue is figuring out how to represent this information in JSON format, and subsequently handling that data. Makeshift texturing could be achieved by passing an image over a buffer at present, but I believe this is far slower than using the intended API calls.
- **Optimisation** — at present, Palette makes many redundant calls to the WebGL API. While these do not affect a project on the scale of Sketch, for larger scale rendering tasks this could begin to become a problem. This should be possible to mitigate by making use of tools such as WebGL Inspector, learning from their output, and writing optimisations from its findings.
- **Shader Loading from URL** — the current implementation forces users to externally load shader JSON files themselves - likely through XML HTTP Requests. Making startup simpler would be a great aid to future users.
- **Draw Mode Extensibility** — while this is possible at the current moment, it requires direct modification of Palette's source due to the current internal architecture. Development of a plugin system may be a worthwhile investment in the future to prevent users from modifying the code themselves to add their required functionality.

Overall, implementing Palette has been an extremely rewarding experience both for the nuance involved in writing such a subsystem and for the value it has provided to the development of other Sketch subcomponents. Team members who have had to use it had a small adjustment time, but appeared to be suitably impressed with the simplicity it provided them.

6.6 Module Facade

Chapter 7

Development Environment

Chapter 8

Online IDE

As part of our overall project plan, we decided that an environment for users to play around with Sketch was an essential deliverable for us to provide - both for our own testing purposes and to offer an interactive sandbox that could show off our project's capabilities.

8.1 Design Process

Our team had not originally considered the production of a development environment for our language, until the topic had been raised by our supervisor in the language development stages. Up until that point we had merely assumed that users would produce programs on a desktop computer in their text editor of choice, before linking the code up to the module on a web-page via the JavaScript API. However, mobile devices were left with no means to write their own Sketch programs - this left out a potential user base that we were specifically trying to target by electing to use WebGL.

We began to look at existent online IDEs for various languages to serve as design inspiration for our own. Chief among these was ShaderToy (Figure 8.1a), an online community built around sharing and discussing complex fragment shaders. ShaderToy's design makes the code and the visual output the primary focus for users, with both sections taking up relatively equal screen space



(a) An image of ShaderToy's design layout.

(b) An image of 0x10co.de's design layout.

Figure 8.1: Other online examples.


at opposite ends of the screen. It's noted that the code segment takes up slightly more screen space - this is likely due to the far larger size of GLSL fragment shader code. While many elements of its design are bland, we found the structure of the page to be a perfect fit for graphics dependent languages where the canvas takes centre stage. We even found that the bland colour scheme of the site likely helped the key visual elements - the canvas and the code - to stand out to users by reducing visual clutter and noise elsewhere.

The now defunct 0x10co.de (Figure 8.1b), an online community based around the fictional DCPU-16, was another key influence in this regard - it also followed a minimalistic design scheme, placing code and graphical output at the forefront. We decided upon minimalism as our go-to style after our investigation to give the page and language presentation the purity we desired, and chose to go for a 50-50 split as our programs are intended to be relatively small. The key decision to deviate from the light pages in our two exemplars was made in light of modern trends within text editors such as Sublime Text - developers were regularly finding darker environments to be more aesthetically pleasing and less likely to cause eye-strain. As such, we decided to go for a darker web design; not only for these reasons, but also to enhance the appearance of the canvas by creating a stronger contrast.

8.2 Implementation

On further research, it was noted that 0x10co.de and ShaderToy shared a common library - CodeMirror. CodeMirror is a JavaScript library designed to allow embedding fully formattable text entry boxes into web pages, with support for syntax highlighting, theme selection, various plugins and more. After researching it briefly, we were certain we'd found the right tool for the job.

Creating a new syntax colouration template was at first a difficult task - in essence, it was identical to building a parser for Sketch. Luckily, CodeMirror provides a simple module definition plugin, which reduces the task down to the construction of a finite state machine with the help of regular expressions. The task was far easier to tackle this way, and although it lacks the fine-grained control of writing a custom parser the results are very convincing. An image showing this colouring in action is reproduced for Figure 8.2. The state machine is as follows:



```

1 Line a;
2 function init(){
3     a = {0,1};
4 }
5 function render(){
6     draw march(a);
7 }
8 function march(Point p)->Point{
9     vec2 b = {1,1};
10    string c = "A string";
11    return a+b;
12 }
13

```

Figure 8.2: Sketch's syntax colouring, provided by CodeMirror.

```

1 CodeMirror.defineSimpleMode("mlang",{
2     start:[
3         {regex: /"(?:[^\\"|\\\.])*?"/, token: "string"},
4         {regex: /(function) (\s+) ([a-z$][\w$]*)/,
5           token: ["keyword", null, "variable-2"]},
6         {regex: /#.*/ , token: "comment"},
7         {regex: /\s\/\s.*/ , token: "comment"},
8         {regex: /true|false/, token: "atom"},
9         {regex: /0x[a-f\d]+|[-+]?(?:\.\d+|\d+\.\.?*\d*)(?:e[-+]?[0-9]+)?/i,
10          token: "number"},

```

```

11     {regex: /\b(?:[^\s\|\\\|\\\.)*?\b/, token: "variable"},
12     {regex: /\b(?:function|return|if|for|while|else|do
13         |this|draw|clear|width|height)\b/, token: "keyword"},
14     {regex: /\b(?:int|float|bool|string|Line|Point|Polygon|void
15         |color|vec2|vec3|vec4|Circle)\b/, token: "keyword"},
16     {regex: /\b(?:vector\([234]\)\)\b/, token: "keyword"},
17     {regex: /[+*<=>!]+/, token: "operator"},
18     {regex: /[a-z$][\w$]*/, token: "variable"},
19     {regex: /([a-z$][\w$]*) (?(=\s*)\() /, token: "keyword"}
20 ],
21 comment: [
22
23 ],
24 meta: {
25     dontIndentStates: ["comment"],
26     lineComment: "//"
27 }
28 });

```

Listing 8.1: Syntax Highlighting State Machine

The page itself (Figure 8.3) is a very simple arrangement of blocks, a canvas, and a text area which is in turn replaced by CodeMirror. This is then styled by page-wide CSS designed to match our chosen CodeMirror theme - creating the subdued look we were going for.




Figure 8.3: An image of the Sketch online IDE.

8.3 Evaluation

Being able to produce this page in less than two days, we realised that we had definitely made the correct choice in our library research. The page itself has much room to improve, however - there are vast amounts of unused whitespace, currently links to our documentation are non-functional and the page feels very barren, overall. We believe that the value that this aspect of our project provides is far greater than the time cost involved in creating it, thanks to the simple set-up process and the past experience of team members with HTML and CSS as well as its own intrinsic value.

Chapter 9

Test Plan and test suites

Chapter 10

Conclusion

10.1 Contributions

10.1.1 Report

Writing:

- *David Creigh*: Chapter
- *Kristiyan Dimitrov*: Chapter 6.2
- *Gerard Docherty*: Chapter
- *Darren Findlay*: Chapter
- *Kyle Simpson*: Chapter 6.1, 6.5, 6.6, 7, 8

Editing:

- *David Creigh*: Chapter
- *Kristiyan Dimitrov*: Chapter
- *Gerard Docherty*: Chapter
- *Darren Findlay*: Chapter
- *Kyle Simpson*: Chapter

10.1.2 Program

Appendix A

Appendix A: Source Listing

Includes full source listing of compiler